

Cancellation and Surplus Patterns: A 2-Page Cheat Sheet

GOAL

Find exact or approximate frequent elements in a data stream with small memory and one pass. Central idea: maintain a minimal residue via online cancellation or a surplus (vote margin) potential.

TERMS

Residue: result of deleting “opposite” pairs without changing winners. Vote margin s : unpaired count of current candidate. Phase: segment between resets when s returns to 0.

PATTERN 1 – Boyer–Moore Majority ($> n/2$)

State: (candidate c , margin $s \geq 0$)

Update per x : if $s == 0 \rightarrow c := x; s := 1$ else if $x == c \rightarrow s++$ else $\rightarrow s--$

Invariant: online simulation of deleting unequal pairs. After any prefix, residue is empty or k copies of one value; (c, s) equals that residue. If a true majority exists, residue over full stream is non-empty and equals it.

Verify: second pass counts(c) $> n/2$

Merge/distribute: yes; reduce chunk residues with same rule, then verify globally.

Complexity: $O(n)$ time, $O(1)$ space.

Boyer–Moore pseudocode (reference)

```
s=0; for x in stream:
```

```
    if s==0: c=x; s=1
```

```
    elif x==c: s+=1
```

```
    else: s-=1
```

```
# verify c by counting
```

PATTERN 2 – Misra–Gries Heavy Hitters ($> n/k$)

Goal: candidates for all items with freq $> n/k$.

State: up to $k-1$ pairs (item, count).

Update: if x tracked $\rightarrow ++$; else if room $\rightarrow \text{insert}(x, 1)$; else decrement all counts and delete zeros.

Guarantee: every true heavy hitter appears among candidates; counts are underestimates by at most total decrements. Verify in second pass.

Notes: exact candidate set with $O(k)$ space and one pass; mergeable across partitions.

PATTERN 3 – Space-Saving (approx top- k with errors)

State: k slots (item, count, error).

Update: hit $\rightarrow \text{count}++$; miss \rightarrow replace min-count slot with $(x, \text{min}+1, \text{min})$.

Guarantee: $\text{stored_count}(x) - \text{error}(x) \leq \text{true_count}(x) \leq \text{stored_count}(x)$.

Use: fixed memory, good for streaming top- k dashboards.

PATTERN 4 – XOR Parity (Single Number)

State: $a=0$; Update: $a \oplus= x$

Invariant: pair duplicates cancel under XOR; a equals element with odd multiplicity.

Use: “find the unique among pairs/triples” variants (extend with bitwise FSM for triples).

PATTERN 5 – Monotone Stack/Deque (Next Greater, Stock Span, Sliding Window Min)

State: stack/queue keeping monotonicity.

Update: pop while invariant violated; push x .

Invariant: structure holds only candidates that can survive future dominance checks.

Use: $O(n)$ time range extrema and spans.

POTENTIAL / RESET FAMILY

PATTERN 6 – Kadane (Max Subarray)

State: `best_suffix`, `best`. Update: `best_suffix = max(x, best_suffix + x)`; `best = max(best, best_suffix)`.
Idea: forget history when debt is positive. Merge via (`sum`, `best`, `best_prefix`, `best_suffix`).

PATTERN 7 – Gas Station (Circular Tour)

State: `start`, `tank`, `total`. Update: `tank += gas[i] - cost[i]`; if `tank < 0`: `start = i + 1`; `tank = 0`.
Invariant: last reset wins if `total >= 0`. One pass, $O(1)$.

SKETCHING / WINDOWED

PATTERN 8 – Count-Min Sketch (ϵ, δ approximate counts)

State: `d` hash rows \times `w` counters.
Update: inc one cell per row at `h_j(x)`. Query: min over rows.
Error: overestimates by $\leq \epsilon N$ with prob $\geq 1 - \delta$. Merge: add counters.

PATTERN 9 – DGIM / Smooth Histograms (Sliding Windows)

State: timestamped buckets with exponentially scaled sizes; keep ≤ 2 per size.
Update: on each 1, create bucket; merge oldest pair when > 2 . Query by summing buckets.
Use: approximate counts in last `W` items.

COMPOSITION / QUERIES

PATTERN 10 – Segment Tree with Boyer-Moore Monoid

Node state: (`candidate`, `margin`). Merge: apply Boyer-Moore's cancel/accumulate to children.
Use: range majority queries; verify candidates on demand.

WEIGHTED VARIANTS

PATTERN 11 – Weighted Boyer-Moore

State: (`c`, `s`). For item with weight `w`: if `x == c` \rightarrow `s += w` else `s -= w`; if `s \leq 0` set `c := x` and `s := -s` (or pick next with residual).
Use: majority under costs/weights; verify by weighted count.

SAMPLING

PATTERN 12 – Reservoir Sampling (`k` items)

State: sample `S`, `t`. Update: `t += 1`; with prob `k/t` replace a random element in `S` with `x`.
Invariant: each seen item in `S` with prob `k/t`. Merge: weighted reservoirs.

DESIGN CHECKLIST

- 1) Exact vs approximate? Threshold known ($> n/2$, $> n/k$)?
- 2) Stream constraints: one pass, memory cap, verification possible?
- 3) Mergeability: per-partition summaries must compose.
- 4) Windowed vs global? If windowed, prefer deque or DGIM-style summaries.
- 5) Adversarial order? Keep proofs independent of input order.
- 6) Weighted inputs? Use weighted cancellation or reweighting.

CHOOSE (RULE OF THUMB)

- Exact global majority: Boyer-Moore + verify.
- All exact $> n/k$: Misra-Gries + verify.
- Top-`k` under tight RAM, allow error: Space-Saving or Count-Min + heap.
- Unique by parity: XOR or bit-FSM.
- Range queries: segment tree with Boyer-Moore monoid + verify.
- Windowed extrema: monotone deque; windowed counts: DGIM.
- Running "best with resets": Kadane, Gas-station.

PITFALLS

- Skipping verification when majority may not exist.
- Misinterpreting counts from Misra-Gries as exact.
- Using sketches where deletions or windows are required (needs specialized variants).
- Ignoring merge semantics in distributed settings.
- Allowing negative margins without reset in weighted variants.

MINIMAL REFERENCES

- Boyer & Moore, 1981. "MJRTY – A Fast Majority Vote Algorithm."
- Misra & Gries, 1982. "Finding Repeated Elements."
- Metwally et al., 2005. "Efficient Computation of Frequent and Top-`k` Elements..."
- Cormode & Muthukrishnan, 2005. "An Improved Data Stream Summary: the Count-Min Sketch."